

Scoring & Evaluation

Your algorithm is not evaluated on absolute profit alone. We utilize a risk-adjusted utility function designed to reward consistent returns and heavily penalize volatility.

The Objective Function

Your final score is calculated using the following piecewise function:

$$\text{Score} = \begin{cases} \mu \cdot \frac{\text{SR}^2}{\text{SR}^2 + 1}, & \mu \geq 0 \text{ and } \sigma \geq 10^{-10} \\ \mu, & \mu < 0 \text{ or } \sigma < 10^{-10} \end{cases}$$

Where:

- $\mu = \text{Mean(PL)}$
- $\sigma = \text{StdDev(PL)}$
- $\text{SR} = \sqrt{250} \times \frac{\mu}{\sigma}$

How it works: If your strategy loses money on average ($\mu < 0$), your score is simply your mean daily loss. If your strategy is profitable ($\mu \geq 0$) with meaningful variance in daily PL, your mean profit is scaled by a factor of $\frac{\text{SR}^2}{\text{SR}^2 + 1}$. Strategies with high returns but massive volatility (low SR) will see their scores severely discounted. If your daily PL has almost no variance ($\sigma < 10^{-10}$, e.g. a near-constant tiny gain every day), the Sharpe scaling is skipped and your score is just μ directly - this mainly matters for strategies sitting right at the [minimum trading activity](#) threshold.

Judging Criteria

In the final round, your performance is assessed on both quantitative results and qualitative methodology:

- **Quantitative Performance (50%):** Assessed using the Objective Function above against unseen, out-of-sample data.
- **Technical Presentation (50%):** Finalists present their methodology to a panel of researchers and traders. Judges evaluate:
 - Clarity of technical maturity and strategy logic.

- Quality of communication and presentation style.
- Team cohesion and Q&A responsiveness.

Dataset Release Schedule

The evaluation utilizes a total of **2,000 days** of simulated price data, released in stages to test out-of-sample performance and prevent overfitting. At every stage, the live leaderboard is always scored on days you have **not** been given yet - never on data already sitting in your local `prices.txt` - so the leaderboard reflects real predictive performance rather than hindsight.

Confirmed so far:

Date	Stage	You receive locally	Leaderboard/scoring window	Window size
July 8	Testing Round starts	Days 1–500	Days 501–750 (hidden)	250 days

The exact release structure for the General Round and Finals - including how much data is released at each step and how the leaderboard window moves - is still being finalised. **This section will be updated once that's confirmed**, ahead of the General Round starting. In broad terms: expect further stages to release progressively more historical data while continuing to score you on days you haven't seen yet, working toward the full 2,000-day dataset by the end of Finals.